



# Introduction to Deep Learning

Amrita Vishwa Vidyapeetham  
Amritapuri Campus



# Transformers : Self Attention

---

## Attention Is All You Need

---

Ashish Vaswani\*  
Google Brain  
avaswani@google.com

Noam Shazeer\*  
Google Brain  
noam@google.com

Niki Parmar\*  
Google Research  
nikip@google.com

Jakob Uszkoreit\*  
Google Research  
usz@google.com

Llion Jones\*  
Google Research  
llion@google.com

Aidan N. Gomez\* †  
University of Toronto  
aidan@cs.toronto.edu

Lukasz Kaiser\*  
Google Brain  
lukaszkaizer@google.com

Illia Polosukhin\* ‡  
illia.polosukhin@gmail.com

### Abstract

The dominant sequence transduction models are based on complex recurrent or convolutional neural networks that include an encoder and a decoder. The best performing models also connect the encoder and decoder through an attention mechanism. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely. Experiments on two machine translation tasks show these models to be superior in quality while being more parallelizable and requiring significantly less time to train. Our model achieves 28.4 BLEU on the WMT 2014 English-to-German translation task, improving over the existing best results, including ensembles, by over 2 BLEU. On the WMT 2014 English-to-French translation task, our model establishes a new single-model state-of-the-art BLEU score of 41.8 after training for 3.5 days on eight GPUs, a small fraction of the training costs of the best models from the literature. We show that the Transformer generalizes well to other tasks by applying it successfully to English constituency parsing both with large and limited training data.

- Google Brain , University of Toronto 2017 [Vaswani et al. 2017].

# Why Transformers?

- Recurrent architectures rely on sequential processing of input at the encoding step that results in computational inefficiency, as the processing cannot be parallelized
- Transformer architecture completely eliminates sequential processing and recurrent connections. It relies only on **self attention mechanism** to capture global dependencies between input and output
- Significant **parallel processing, shorter training time** and **higher accuracy** for Machine Translation without any recurrent component

## Challenges of RNN

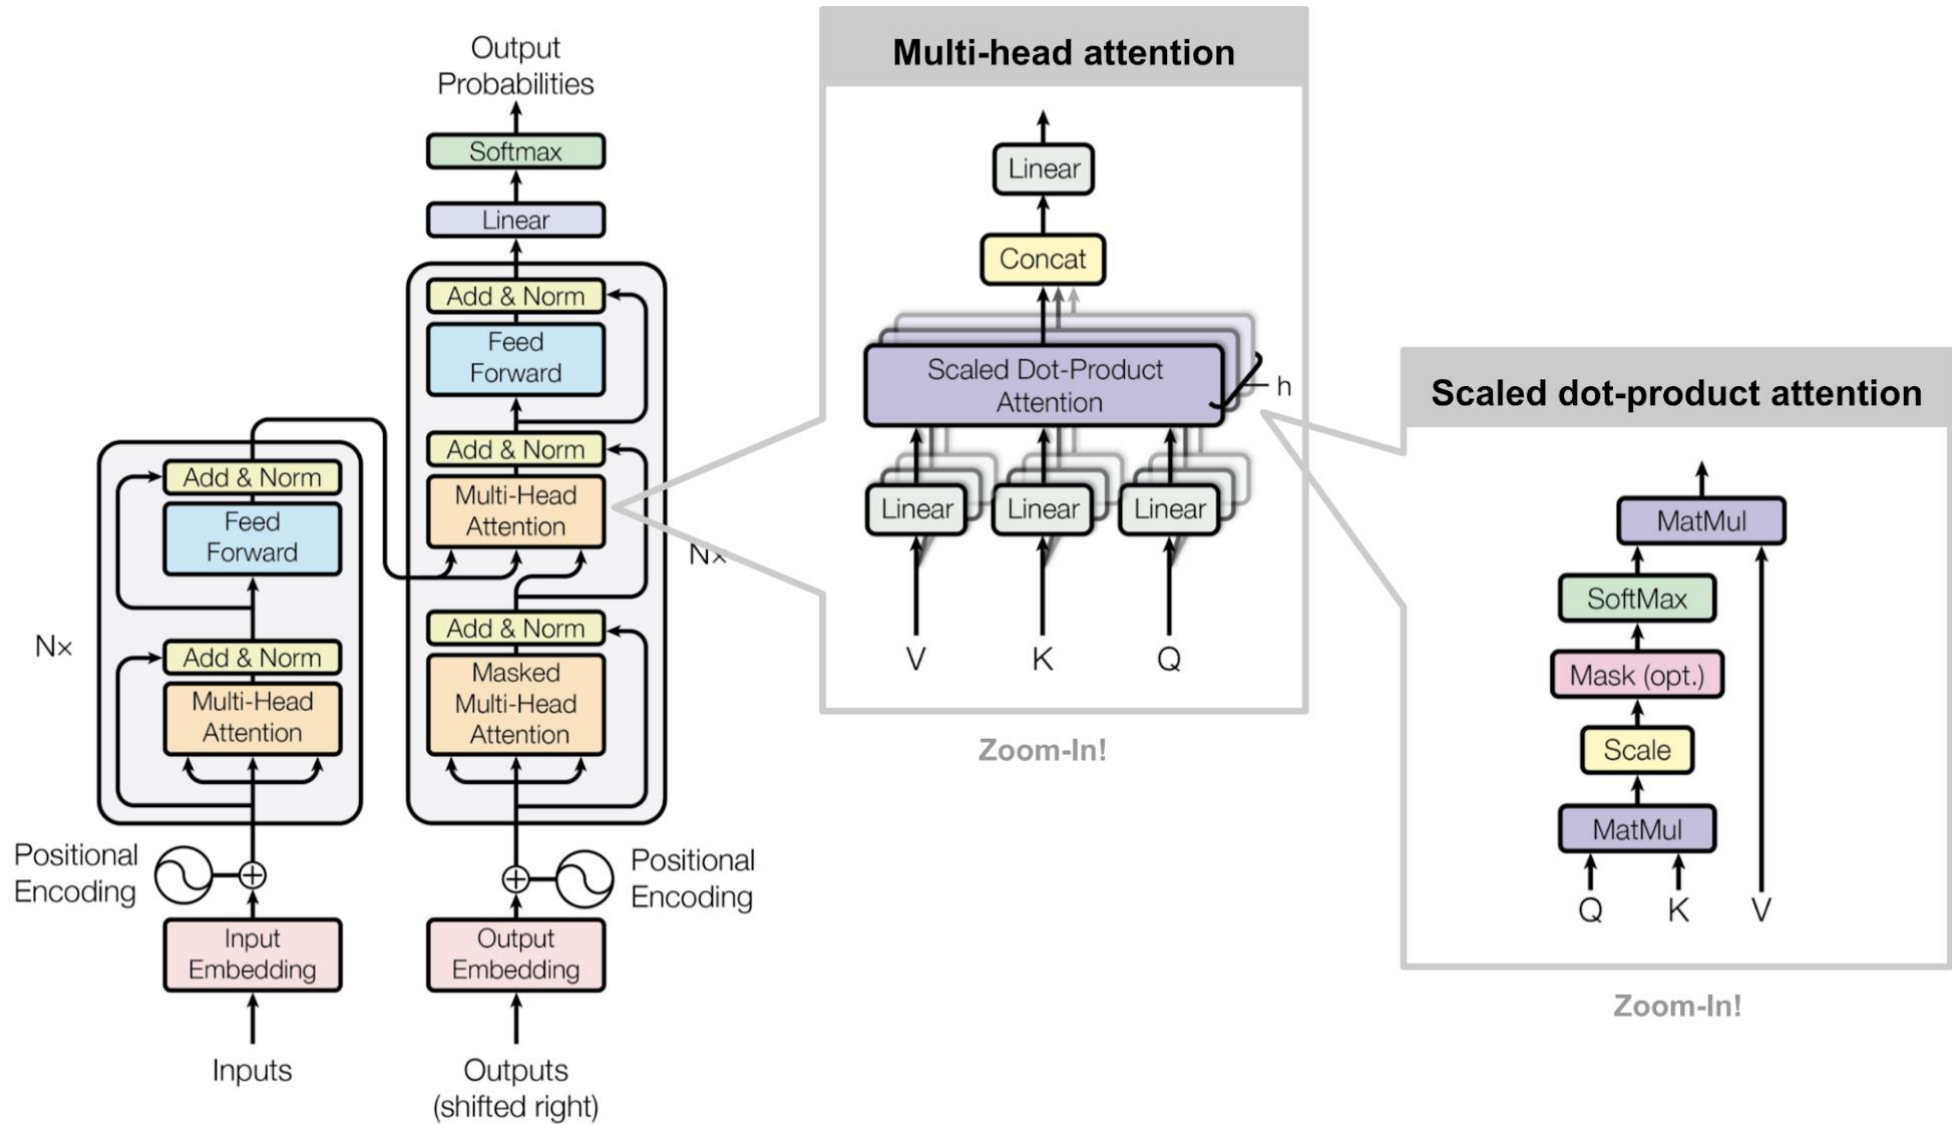
- Long Range Dependencies
- Gradient vanishing and Exploding
- Large # of training steps
- Recurrence prevents parallel computation

## Transformer Networks

- Facilitate long range dependencies
- No gradient vanishing and explosion
- Fewer training steps
- No recurrence and that facilitate parallel computation



# Transformers




# Transformers – Machine Translation example



In a machine translation application, it would take a sentence in one language, and output its translation in another.

Courtesy : [jalammar.github.io](https://github.com/jalammar)

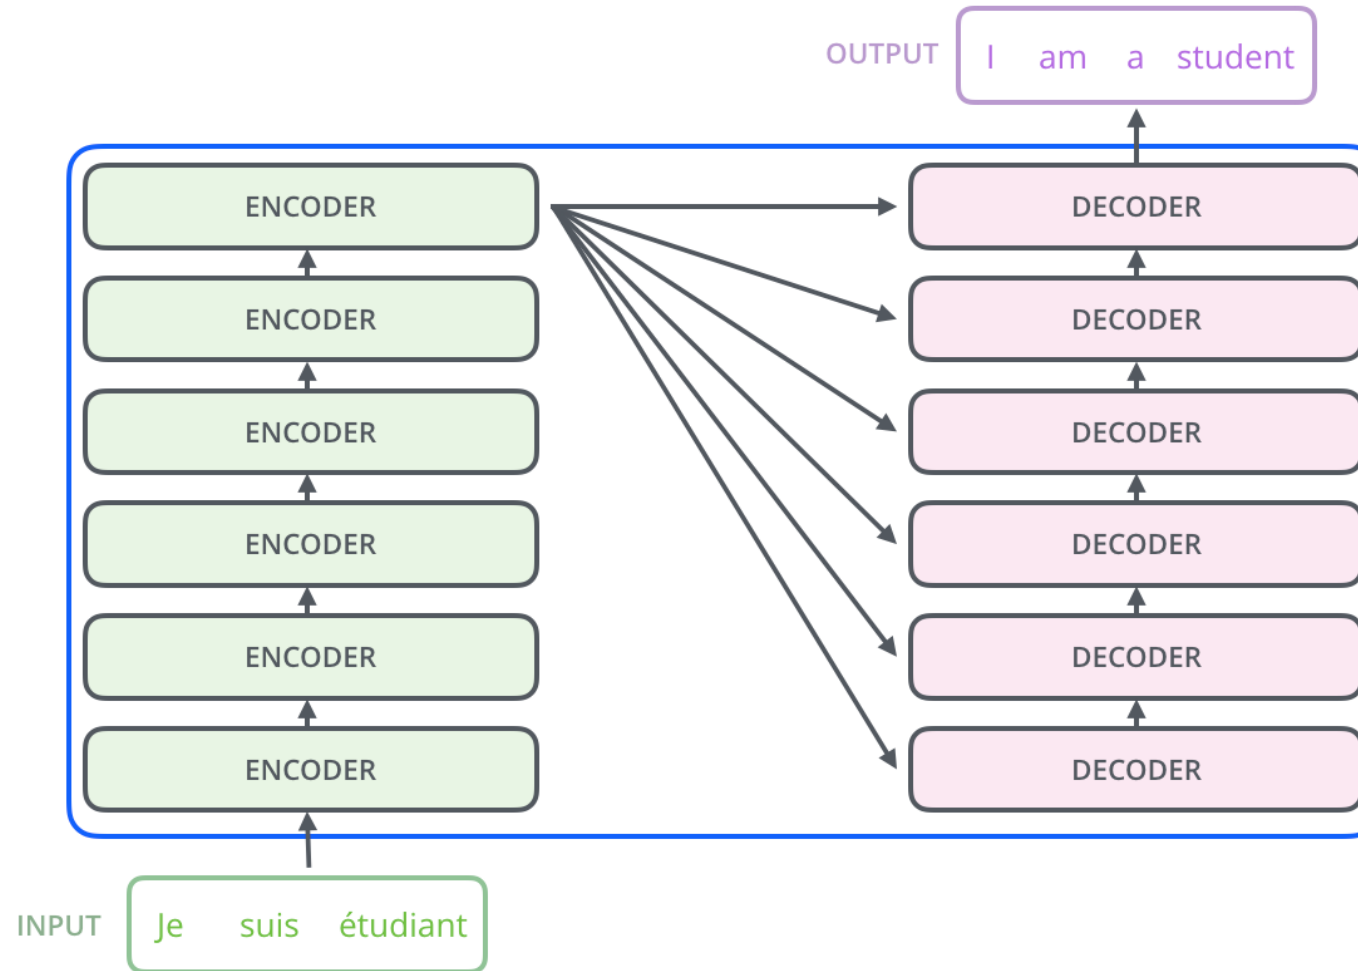
# Transformers – Machine Translation example



An encoding component, a decoding component, and connections between them.

<http://jalammar.github.io>

# Transformers – Machine Translation example




The encoders are all identical in structure (yet they do not share weights).

Each one is broken down into two sub-layers:

The encoding component is a stack of encoders. The decoding component is a stack of decoders of the same number.



# Transformers – Machine Translation example



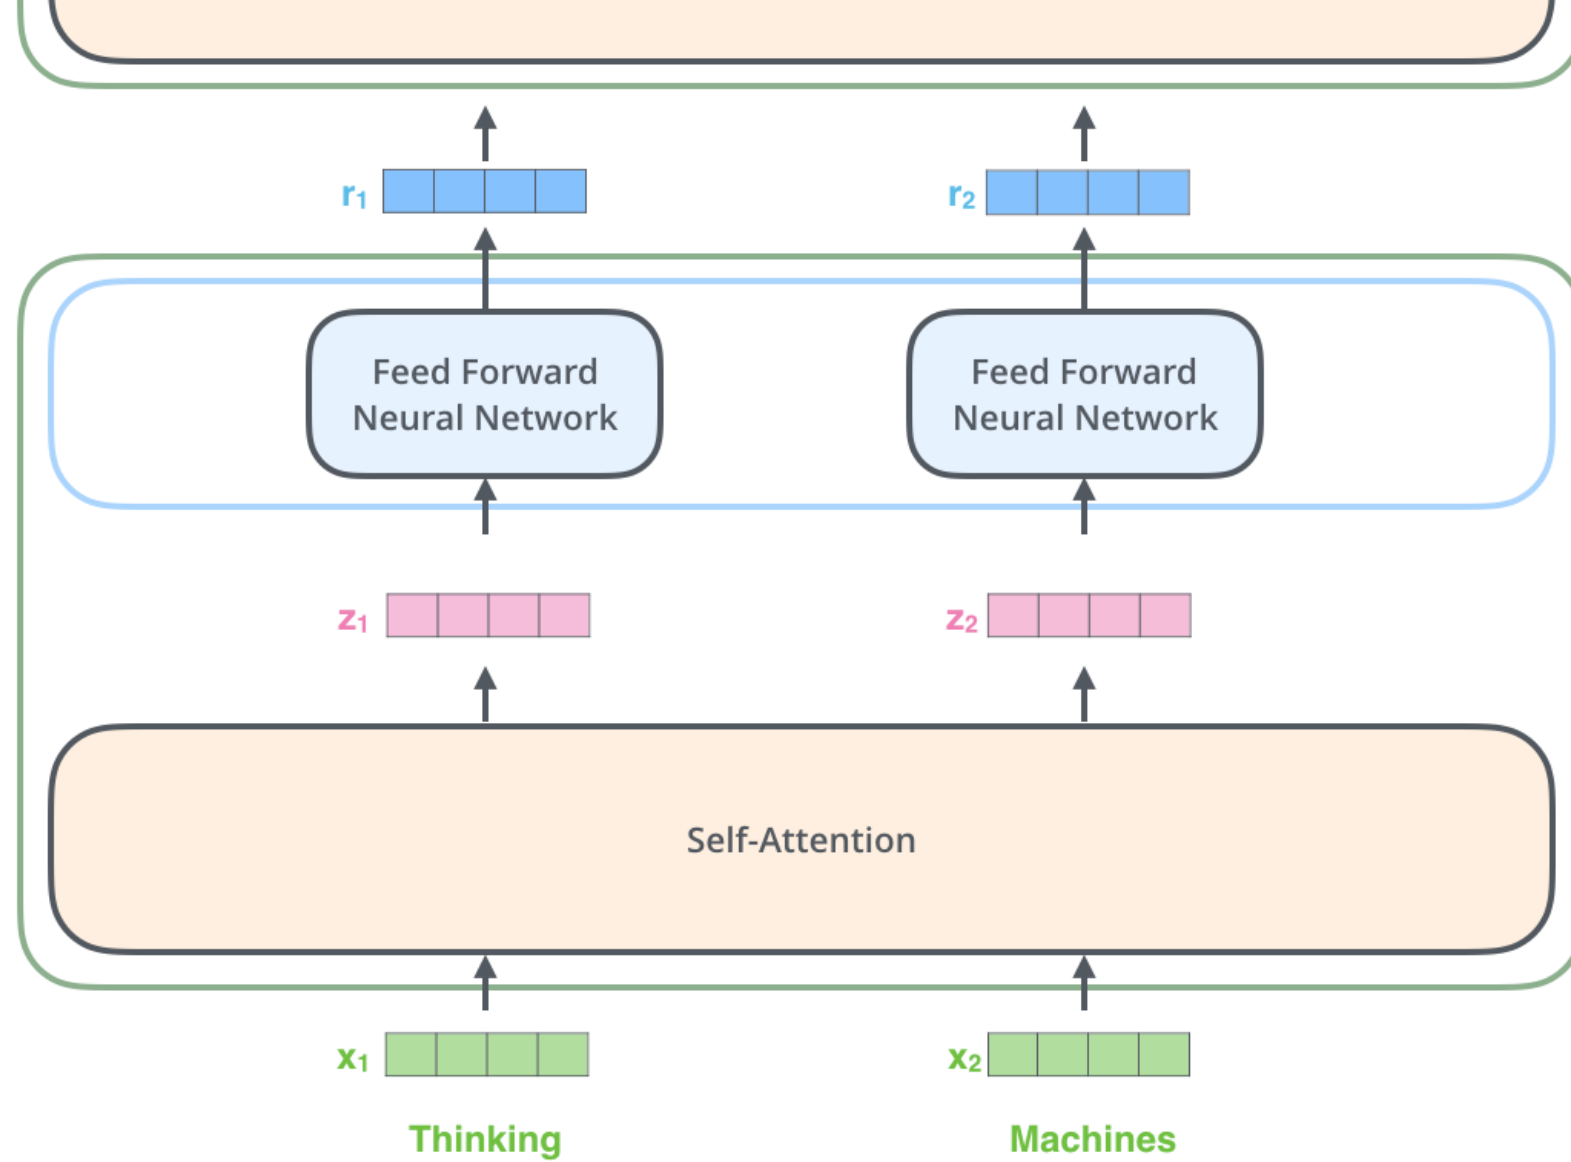
The encoder's inputs first flow through a self-attention layer – a layer that helps the encoder look at other words in the input sentence as it encodes a specific word.

The outputs of the self-attention layer are fed to a feed-forward neural network.

<http://jalammar.github.io>

ENCODER #2

ENCODER #1



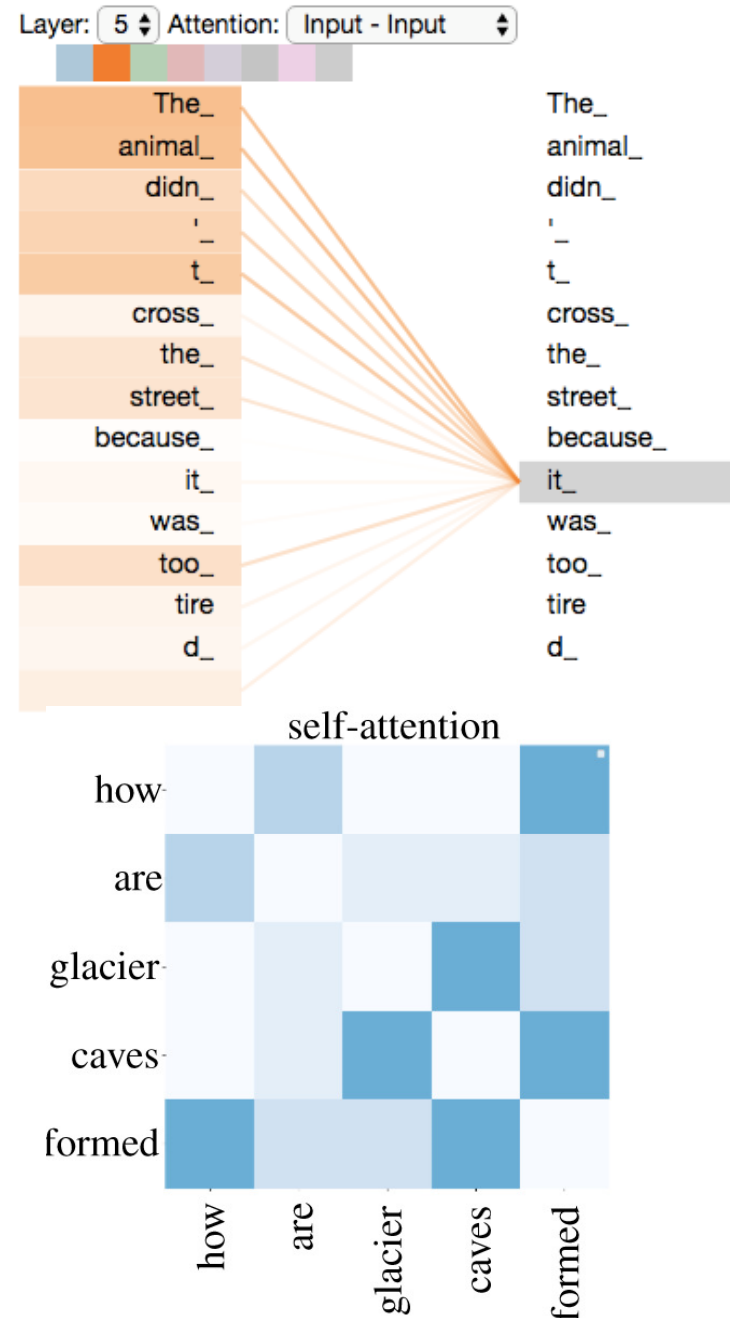
# Self-Attention

The animal didn't cross the street because it was too tired

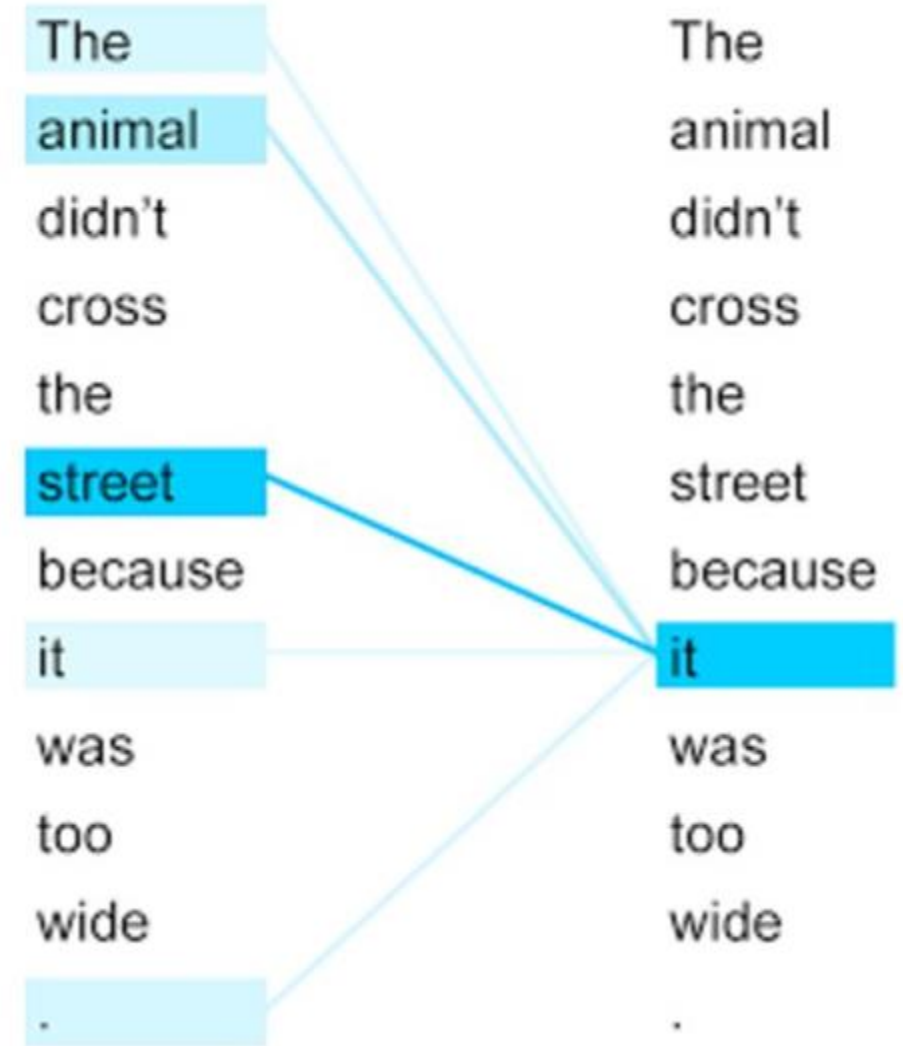
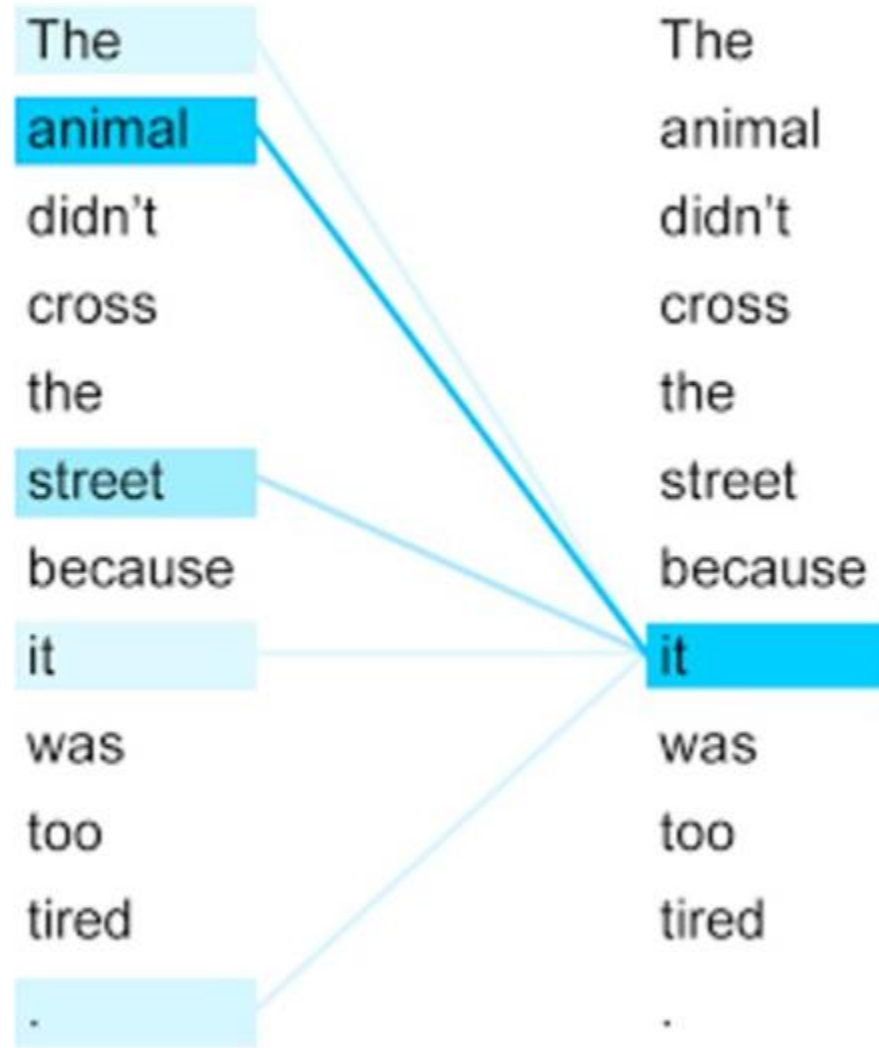
What does “it” in this sentence refer to? Is it referring to the *street* or to the *animal*? It’s a simple question to a human, but not as simple to an algorithm.

When the model is processing the word “it”, self-attention allows it to associate “it” with “animal”.

Self-attention is the method the Transformer uses to bake the “understanding” of other relevant words into the one we’re currently processing.



# Self attention



# Self Attention

- **Query vector (  $q_i$  )**
  - what we are looking for in the sequence
- **Key vector (  $k_i$  )**
  - what the element is “offering”,
- **Value vector (  $v_i$  )**
  - This feature vector is the one we want to average over
- **Score**
  - To rate which elements we want to pay attention to

youtube.com/results?search\_query=Amrita+Vishwa+Vidyapeetham+Amritapuri

Bookmarks IEEE Xplore - Home iGoogle Lenovo Recommen... Dashboard Breast cancer predi... What is better, k-ne... Logistic Regression... Downloads

YouTube

Home Explore Shorts Subscriptions Library History Your videos Watch later Liked videos Attention MEchanism...

SUBSCRIPTIONS Shusen Wang Google Cloud Tech You do it Anon Researcher Dr. Interior Valerio Velardo - The ... Yannic Kilcher Show 56 more

Query → Amrita Vishwa Vidyapeetham Amritapuri

key1

key2

key3

key4

FILTERS

Amrita Vishwa Vidyapeetham, Amritapuri Campus, Kollam, Kerala  
152K views • 2 years ago  
Amrita Vishwa Vidyapeetham  
The Amritapuri campus, nestled in the backdrop of the beautiful village of Vallikavu, provides warmth

Amrita Vishwa Vidyapeetham, Amritapuri Campus  
17.4K subscribers • 223 videos  
This channel is to publicize the programmes offered at Amrita University, Amritapuri Campus.

Amritavarsham 61  
33K views • 7 years ago  
Amrita Vishwa Vidyapeetham, Amritapuri Campus  
Ammma's 61st Birthday Celebration @ Amritapuri.

Unity in Diversity , Amrita University, Amritapuri  
22K views • 6 years ago  
Amrita Vishwa Vidyapeetham, Amritapuri Campus  
Unity in Diversity , Amrita University, Amritapuri.



# Steps in Self Attention

- Query vector (  $q_i$  )
- Key vector (  $k_i$  )
- Value vector (  $v_i$  )

The screenshot shows a YouTube search results page for the query "Amrita Vishwa Vidyapeetham Amritapuri". The search bar at the top contains the query, with a red arrow pointing to it labeled "Query". On the left side, a red bracket labeled "key1" points to the search results. The first result is a video titled "Amrita Vishwa Vidyapeetham, Amritapuri Campus, Kollam, Kerala" with 152K views and 2 years ago. A red arrow points to the video thumbnail, which is labeled "Maximum match, highest attention". Below the video, the text "Extract Values based on highest attention" is written in red.

# Steps in Self Attention –

## First step

Create 3 vectors

- **Query vector (  $q_i$  )**
- **Key vector (  $k_i$  )**
- **Value vector (  $v_i$  )**

From each of the encoder's input vectors –ie the embedding of each word).

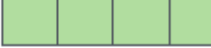
These 3 vectors are created by multiplying the embedding by three matrices that we trained during the training process

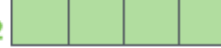
Input

Thinking

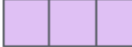
Machines

Embedding


$X_1$  

$X_2$  

Queries

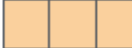
$q_1$  

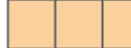
$q_2$  



$W^Q$

Keys

$k_1$  

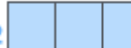
$k_2$  



$W^K$

Values

$v_1$  

$v_2$  



$W^V$


# Steps in Self Attention

## The second step

- Calculate a score : dot product of the **query vector** with the **key vector** of the respective word

$$q_1 \cdot k_1 = 112$$

$$q_1 \cdot k_2 = 96$$



Input

Embedding

Queries

Keys

Values

Score

Divide by 8 ( $\sqrt{d_k}$ )

Softmax

Softmax  
X  
Value

Sum

Thinking

Machines

$x_1$  [ ] [ ] [ ] [ ]

$x_2$  [ ] [ ] [ ] [ ]

$q_1$  [ ] [ ] [ ]

$q_2$  [ ] [ ] [ ]

$k_1$  [ ] [ ] [ ]

$k_2$  [ ] [ ] [ ]

$v_1$  [ ] [ ] [ ]

$v_2$  [ ] [ ] [ ]

$$q_1 \cdot k_1 = 112$$

$$q_1 \cdot k_2 = 96$$

14

12

0.88

0.12

$v_1$  [ ] [ ] [ ]

$v_2$  [ ] [ ] [ ]

$z_1$  [ ] [ ] [ ]

$z_2$  [ ] [ ] [ ]

# Steps in Self Attention

## Third step

Word	q vector	k vector	v vector	score	score / 8
thinking	$q_1$	$k_1$	$v_1$	$q_1 \cdot k_1$	$q_1 \cdot k_1 / 8$
Machines		$k_2$	$v_2$	$q_1 \cdot k_2$	$q_1 \cdot k_2 / 8$

## Divide by 8 ( $\sqrt{d_k}$ )

Normalizing : This leads to having more stable gradients

The square root of the dimension of the key vectors used in the paper – 64.

Input

Embedding

Queries

Keys

Values

Score

Divide by 8 (  $\sqrt{d_k}$  )

Softmax

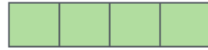
Softmax

X

Value

Sum

Thinking

$x_1$  

$q_1$  

$k_1$  

$v_1$  

$q_1 \cdot k_1 = 112$

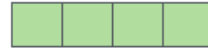
14

0.88

$v_1$  

$z_1$  

Machines

$x_2$  

$q_2$  

$k_2$  

$v_2$  

$q_1 \cdot k_2 = 96$

12

0.12

$v_2$  

$z_2$  


# Steps in Self Attention –

## Fourth step

Word	q vector	k vector	v vector	score	score / 8	Softmax	Softmax * v	Sum
thinking	$q_1$	$k_1$	$v_1$	$q_1 \cdot k_1$	$q_1 \cdot k_1 / 8$	$x_{11}$	$x_{11} * v_1$	$z_1$
Machines		$k_2$	$v_2$	$q_1 \cdot k_2$	$q_1 \cdot k_2 / 8$	$x_{12}$	$x_{12} * v_2$	

Multiply each value vector by the softmax score

The intuition here is to keep intact the values of the word(s) we want to focus on, and drown-out irrelevant words






# Steps in Self Attention –

## Fifth step

Word	q vector	k vector	v vector	score	score / 8	Softmax	Softmax * v	Sum
thinking	$q_1$	$k_1$	$v_1$	$q_1 \cdot k_1$	$q_1 \cdot k_1 / 8$	$x_{11}$	$x_{11} * v_1$	$z_1$
Machines		$k_2$	$v_2$	$q_1 \cdot k_2$	$q_1 \cdot k_2 / 8$	$x_{12}$	$x_{12} * v_2$	

Pass the result through a softmax operation.  
 Softmax normalizes the scores so they're all positive and add up to 1.

Sum up the weighted value vectors. This produces the output of the self-attention layer at this position (for the first word).



# Steps in Self Attention –

## Sixth Step

Word	q vector	k vector	v vector	score	score / 8	Softmax	Softmax * v	Sum <sup>#</sup>
thinking		$k_1$	$v_1$	$q_2 \cdot k_1$	$q_2 \cdot k_1 / 8$	$x_{21}$	$x_{21} * v_1$	
Machines	$q_2$	$k_2$	$v_2$	$q_2 \cdot k_2$	$q_2 \cdot k_2 / 8$	$x_{22}$	$x_{22} * v_2$	$z_2$

We need to score each word of the input sentence against this word. The score determines how much focus to place on other parts of the input sentence as we encode a word at a certain position.

Input

Embedding

Queries

Keys

Values

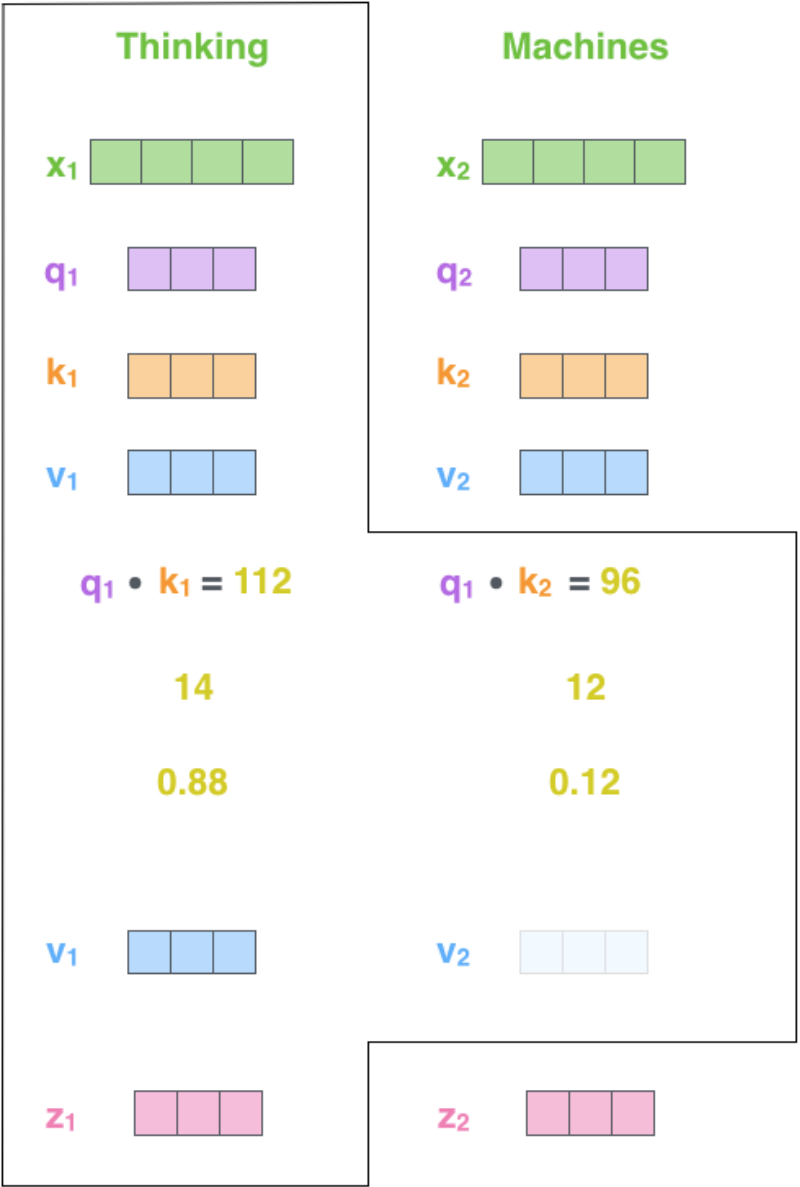
Score

Divide by 8 (  $\sqrt{d_k}$  )

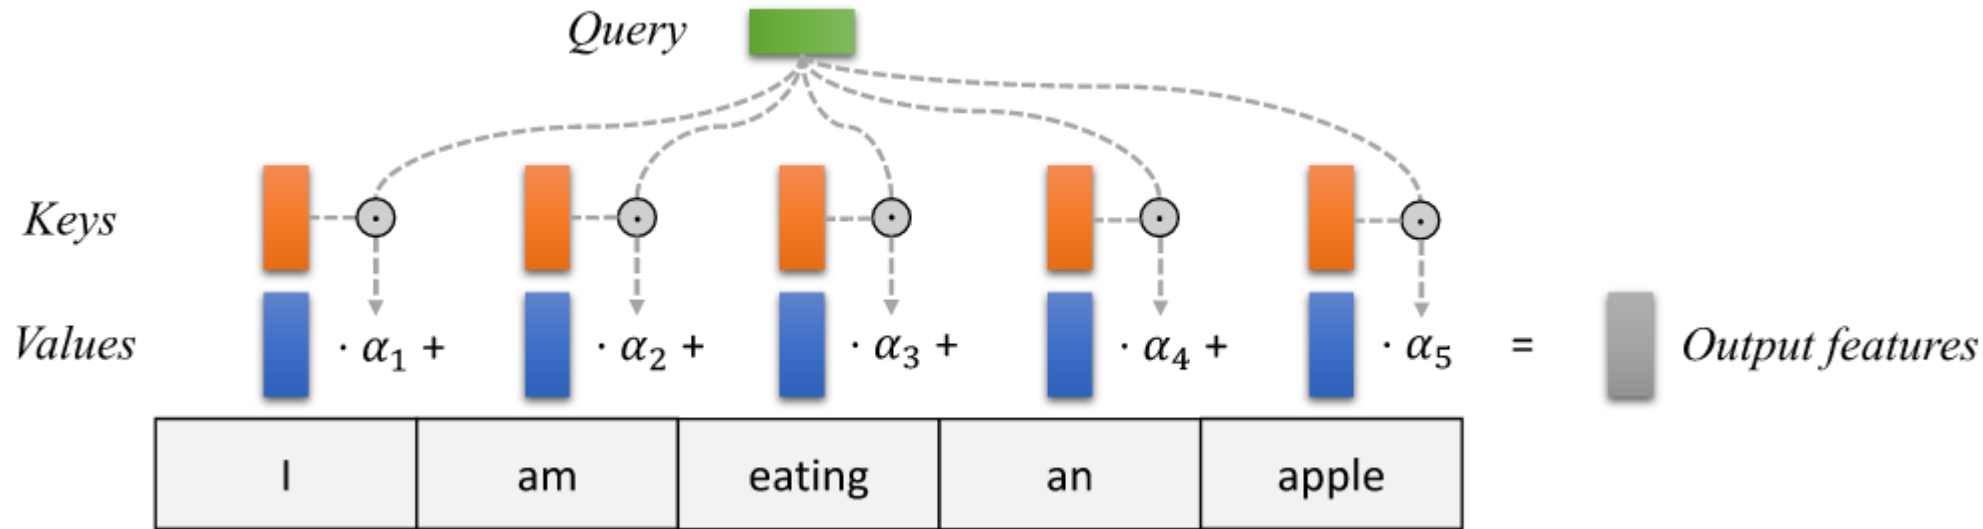
Softmax

Softmax  
X  
Value


Sum



# Self Attention –



Scaled Dot-Product Attention



$$\text{Attention}(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V$$

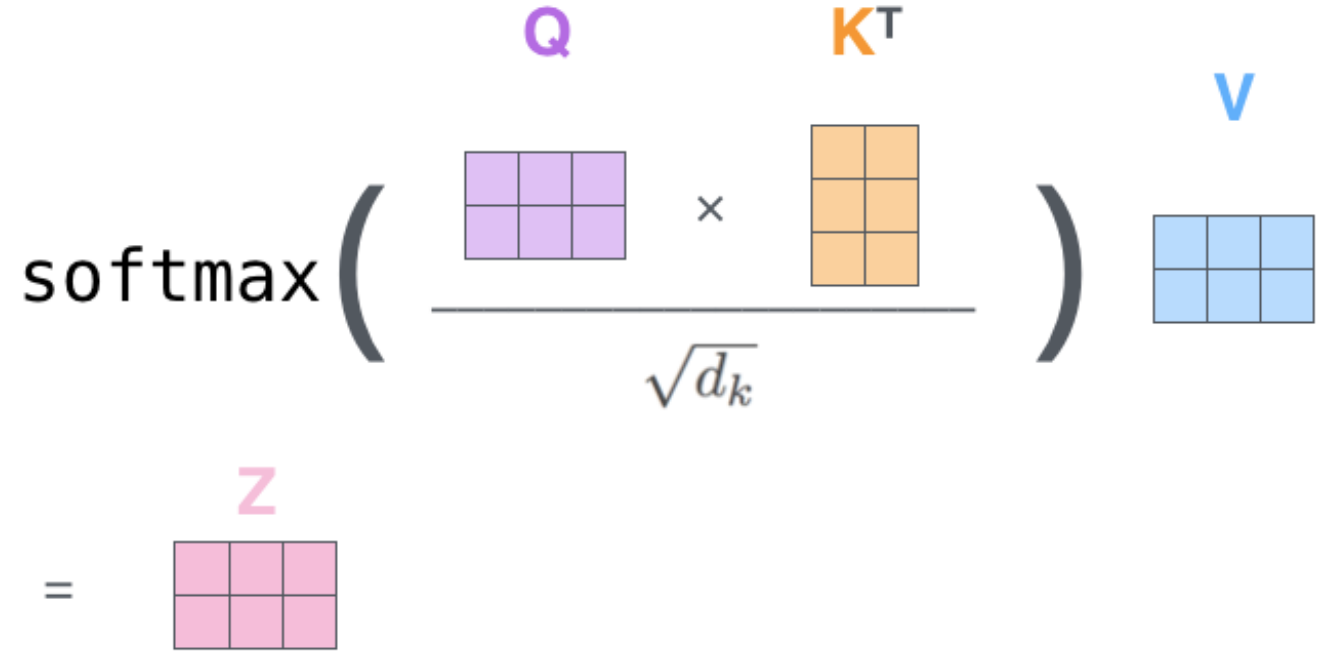
$$\alpha_i = \frac{\exp(f_{\text{attn}}(\text{key}_i, \text{query}))}{\sum_j \exp(f_{\text{attn}}(\text{key}_j, \text{query}))}, \quad \text{out} = \sum_i \alpha_i \cdot \text{value}_i$$

# Matrix Calculation of Self-Attention

$$X \times W^Q = Q$$


$$X \times W^K = K$$


$$X \times W^V = V$$


$$\text{softmax}\left(\frac{Q \times K^T}{\sqrt{d_k}}\right) \times V = Z$$


# “Multi-headed” attention



1) This is our input sentence\*

2) We embed each word\*

3) Split into 8 heads. We multiply  $X$  or  $R$  with weight matrices



4) Calculate attention using the resulting  $Q/K/V$  matrices

5) Concatenate the resulting  $Z$  matrices, then multiply with weight matrix  $W^O$  to produce the output of the layer


Thinking  
Machines





\* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one





...




...



...



Multi-Head Attention

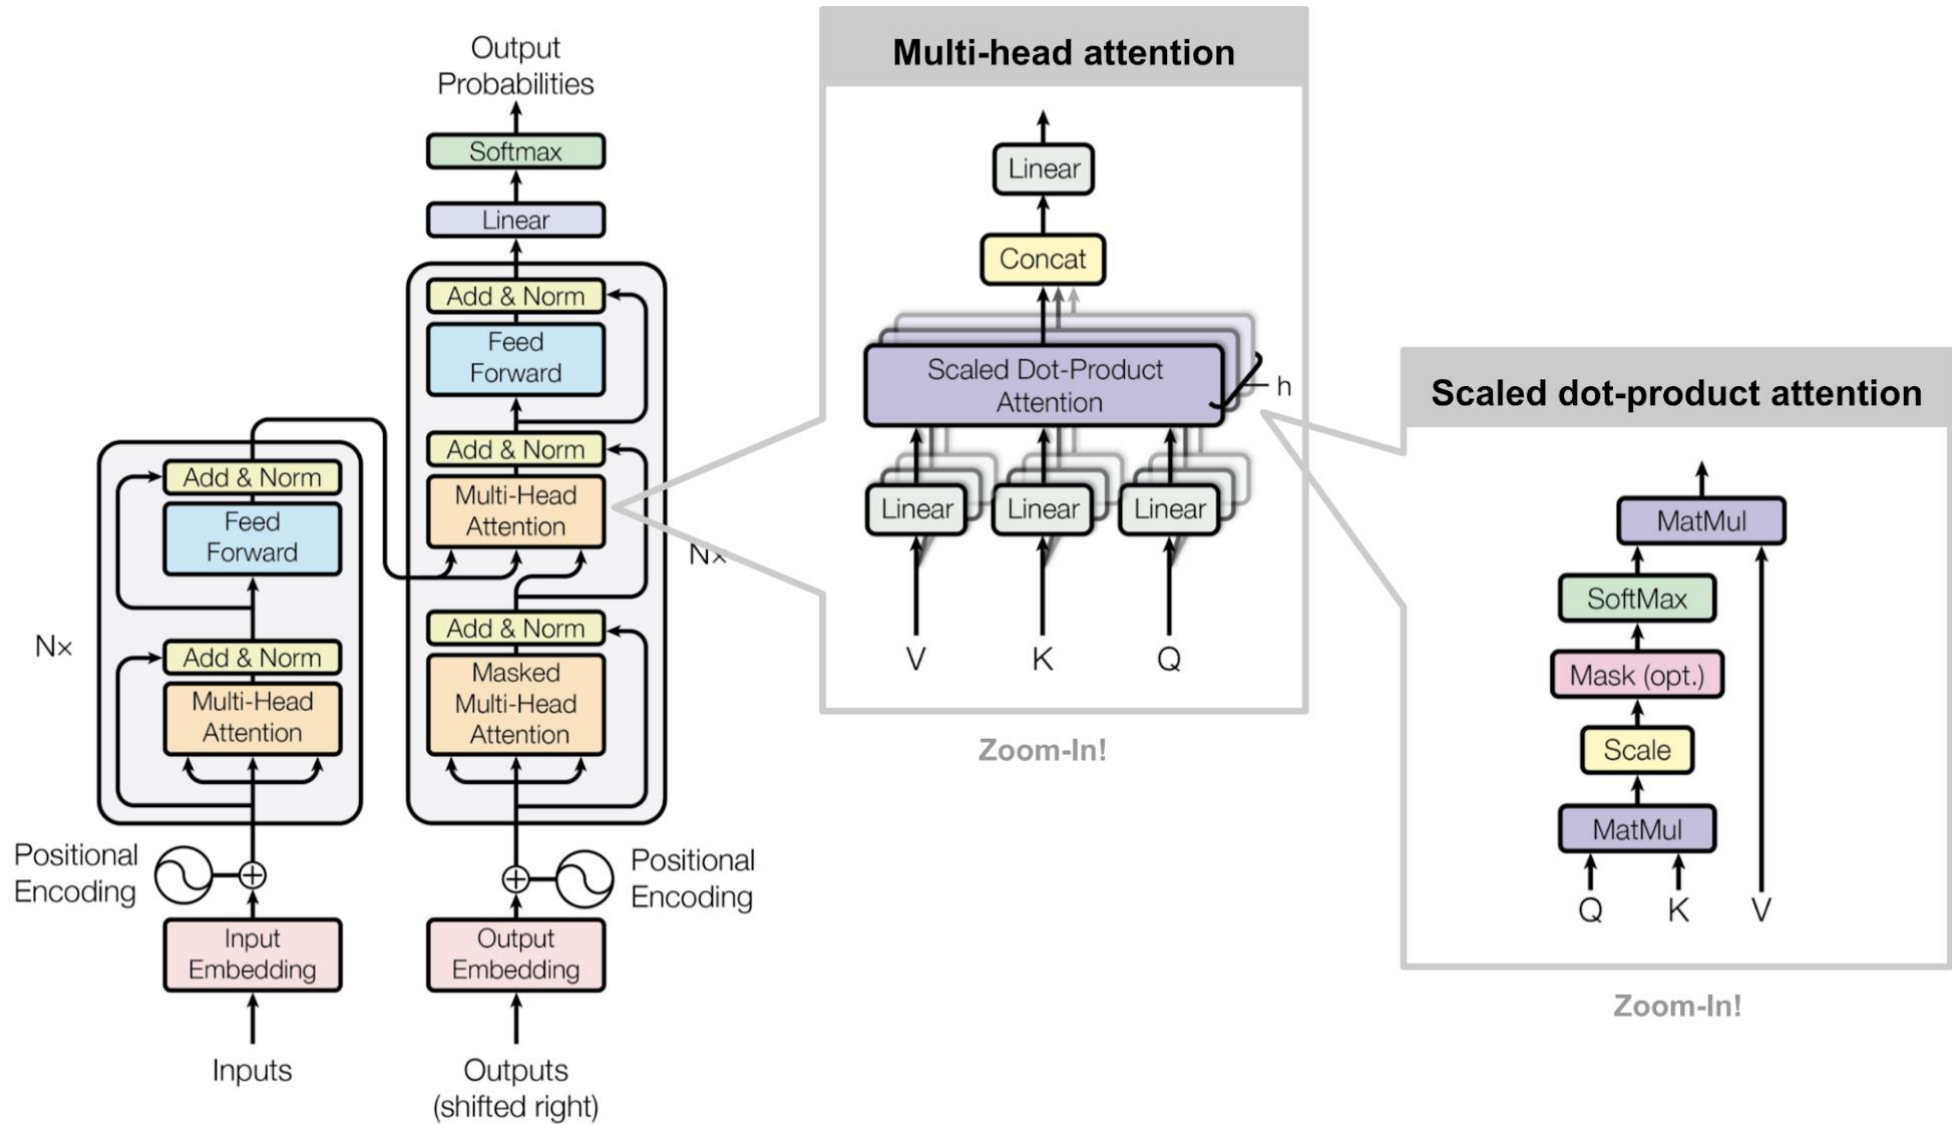


multiple “representation subspaces”

multiple sets of Query/Key/Value weight matrices



# Transformers






# Representing The Order of The Sequence Using Positional Encoding




Multi-Head Attention block is permutation-equivariant, w.r.t its inputs and cannot distinguish whether an input comes before another one in the sequence or not.- hence positional encoding

Permutation-equivariant, :means that if we switch two input elements in the sequence, e.g.  $X_1 \leftrightarrow X_2$  the output is exactly the same besides the elements 1 and 2 switched

Positional encoding



Embedding



Input

Je


suis

étudiant

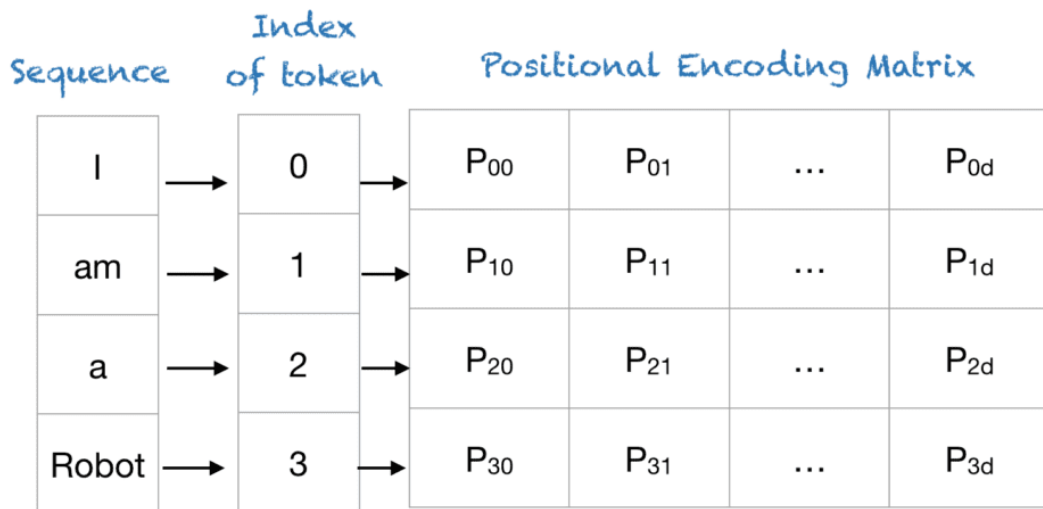
$$PE_{(pos,i)} = \begin{cases} \sin\left(\frac{pos}{10000^{i/d_{\text{model}}}}\right) & \text{if } i \bmod 2 = 0 \\ \cos\left(\frac{pos}{10000^{(i-1)/d_{\text{model}}}}\right) & \text{otherwise} \end{cases}$$

$PE_{(pos,i)}$  represents the position encoding at position pos in the sequence, and hidden dimensionality i.

A real example of positional encoding for 20 words (rows) with an embedding size of 512 (columns)



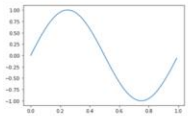
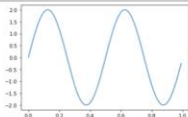
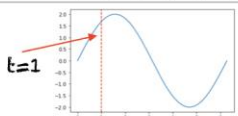
Vaswani et al 2017



## Positional Encoding Layer in Transformers

Positional Encoding Matrix for the sequence 'I am a robot'

Sequence      Index of token,  $k$       Positional Encoding Matrix with  $d=4$ ,  $n=100$

Equation	Graph	Frequency	Wavelength
$\sin(2\pi t)$		1	1
$\sin(2 * 2\pi t)$		2	1/2
$\sin(t)$		$1/2\pi$	$2\pi$
$\sin(ct)$	Depends on $c$	$c/2\pi$	$2\pi/c$

		$i=0$	$i=0$	$i=1$	$i=1$
I	→ 0	$P_{00}=\sin(0)$ = 0	$P_{01}=\cos(0)$ = 1	$P_{02}=\sin(0)$ = 0	$P_{03}=\cos(0)$ = 1
am	→ 1	$P_{10}=\sin(1/1)$ = 0.84	$P_{11}=\cos(1/1)$ = 0.54	$P_{12}=\sin(1/10)$ = 0.10	$P_{13}=\cos(1/10)$ = 1.0
a	→ 2	$P_{20}=\sin(2/1)$ = 0.91	$P_{21}=\cos(2/1)$ = -0.42	$P_{22}=\sin(2/10)$ = 0.20	$P_{23}=\cos(2/10)$ = 0.98
Robot	→ 3	$P_{30}=\sin(3/1)$ = 0.14	$P_{31}=\cos(3/1)$ = -0.99	$P_{32}=\sin(3/10)$ = 0.30	$P_{33}=\cos(3/10)$ = 0.96

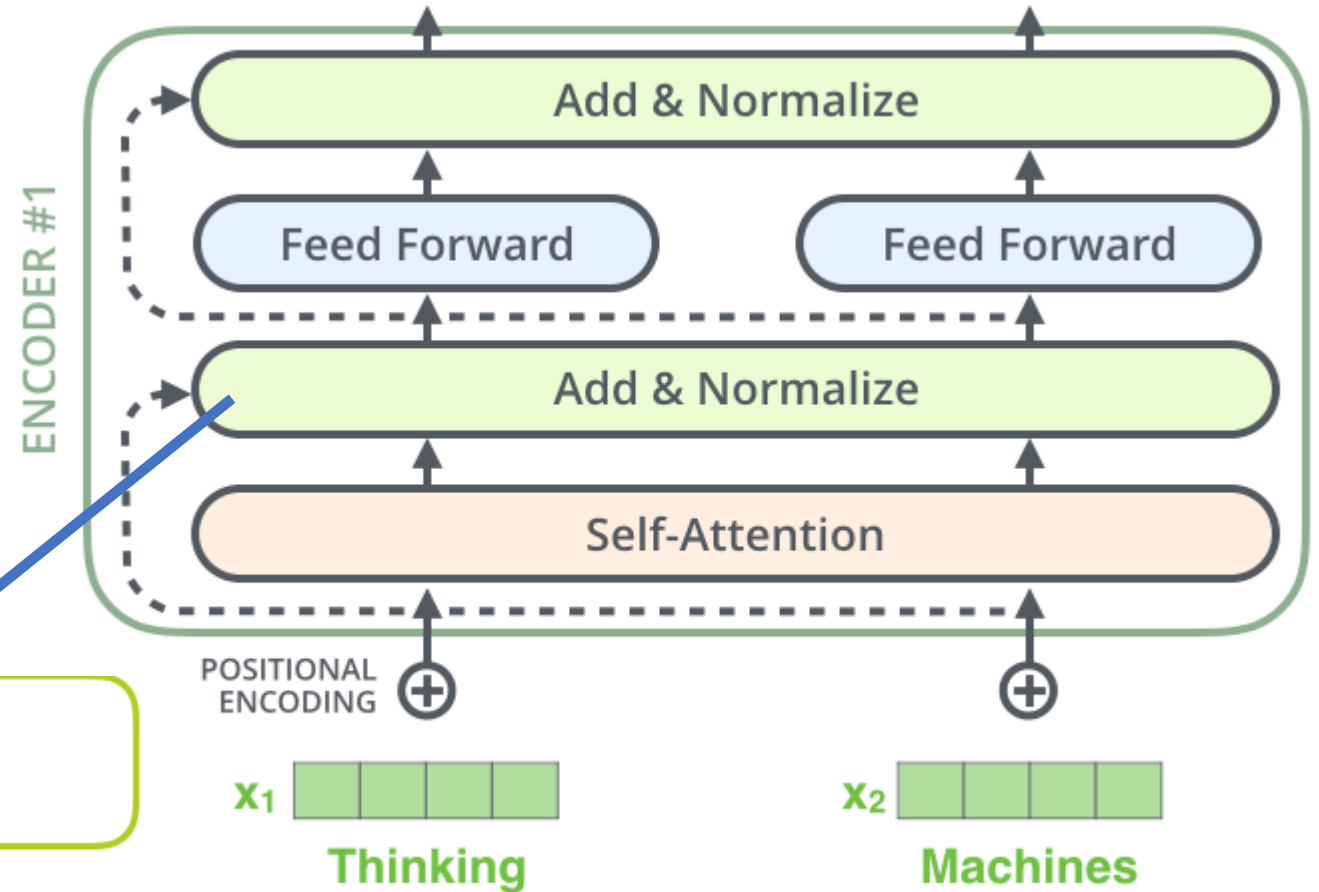

Positional Encoding Matrix for the sequence 'I am a robot'

# Residual connection, Layer Norm, FFN

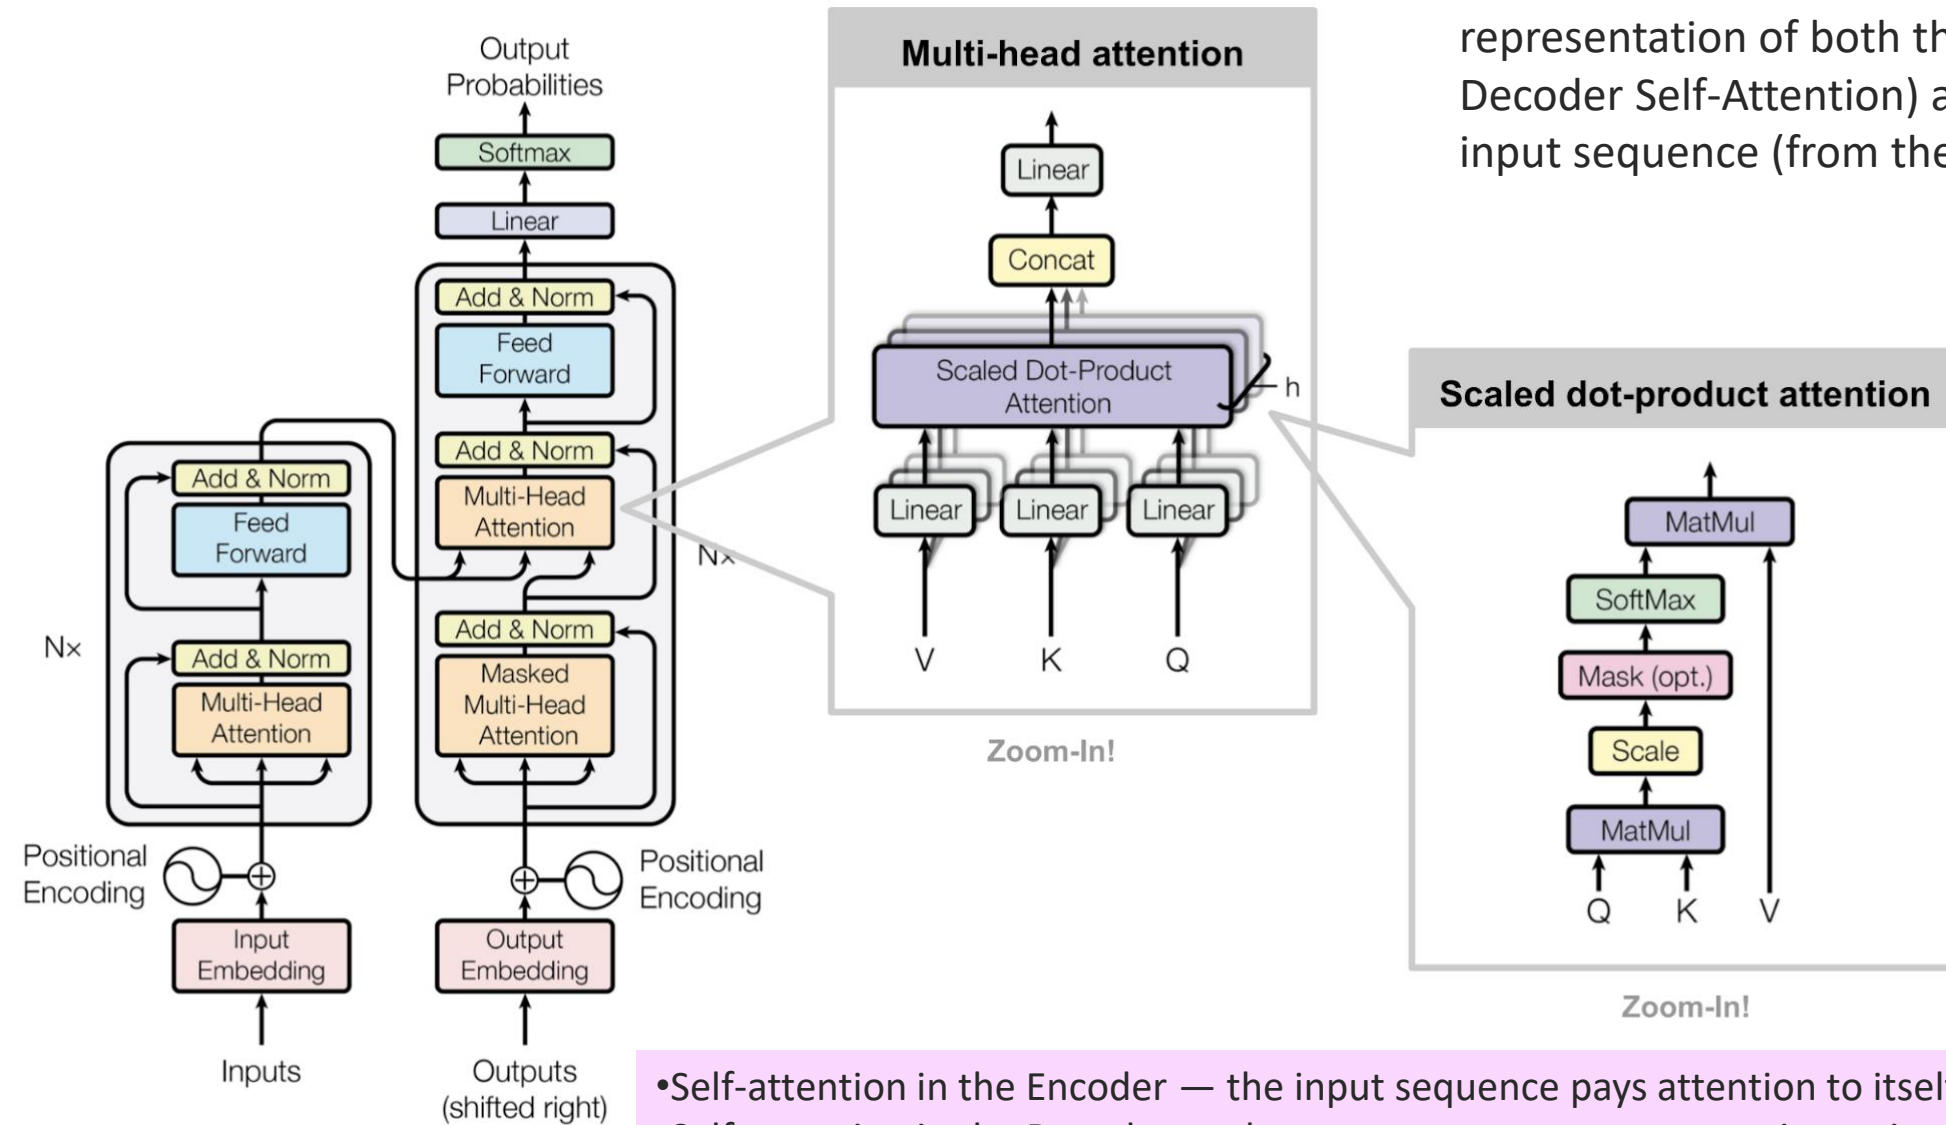
Each encoder has a

- Residual connection around it,
- Followed by a layer normalization step
- Feed forward : -deepens our network, employing linear layers to analyse patterns in the attention layers output.

$$\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$$
$$x = \text{LayerNorm}(x + \text{FFN}(x))$$



# Transformers



The Encoder-Decoder Attention is getting a representation of both the target sequence (from the Decoder Self-Attention) and a representation of the input sequence (from the Encoder stack).

- Self-attention in the Encoder — the input sequence pays attention to itself
- Self-attention in the Decoder — the target sequence pays attention to itself
- Encoder-Decoder-attention in the Decoder — the target sequence pays attention to the input sequence

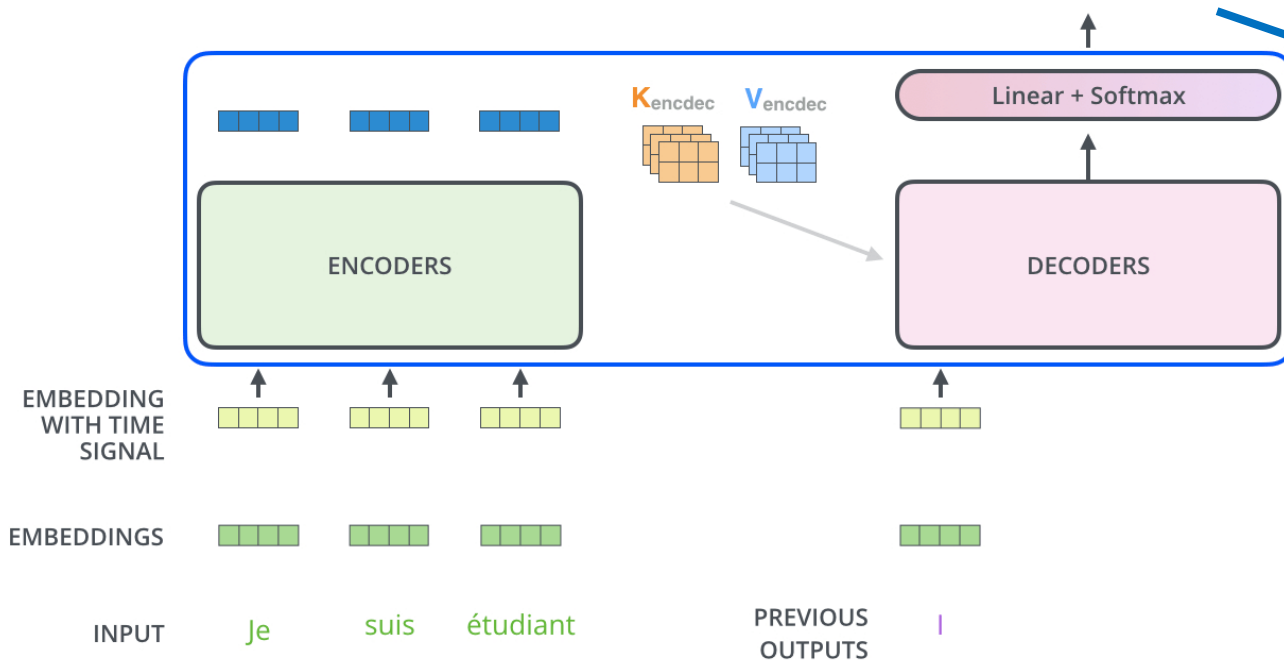


# Transformer

The “Encoder-Decoder Attention” layer works just like multiheaded self-attention, except it creates its Queries matrix from the layer below it, and takes the Keys and Values matrix from the output of the encoder stack.

Decoding time step: 1 2 3 4 5 6

OUTPUT |



Which word in our vocabulary is associated with this index?

Get the index of the cell with the highest value (argmax)


log\_probs

logits

Decoder stack output

am

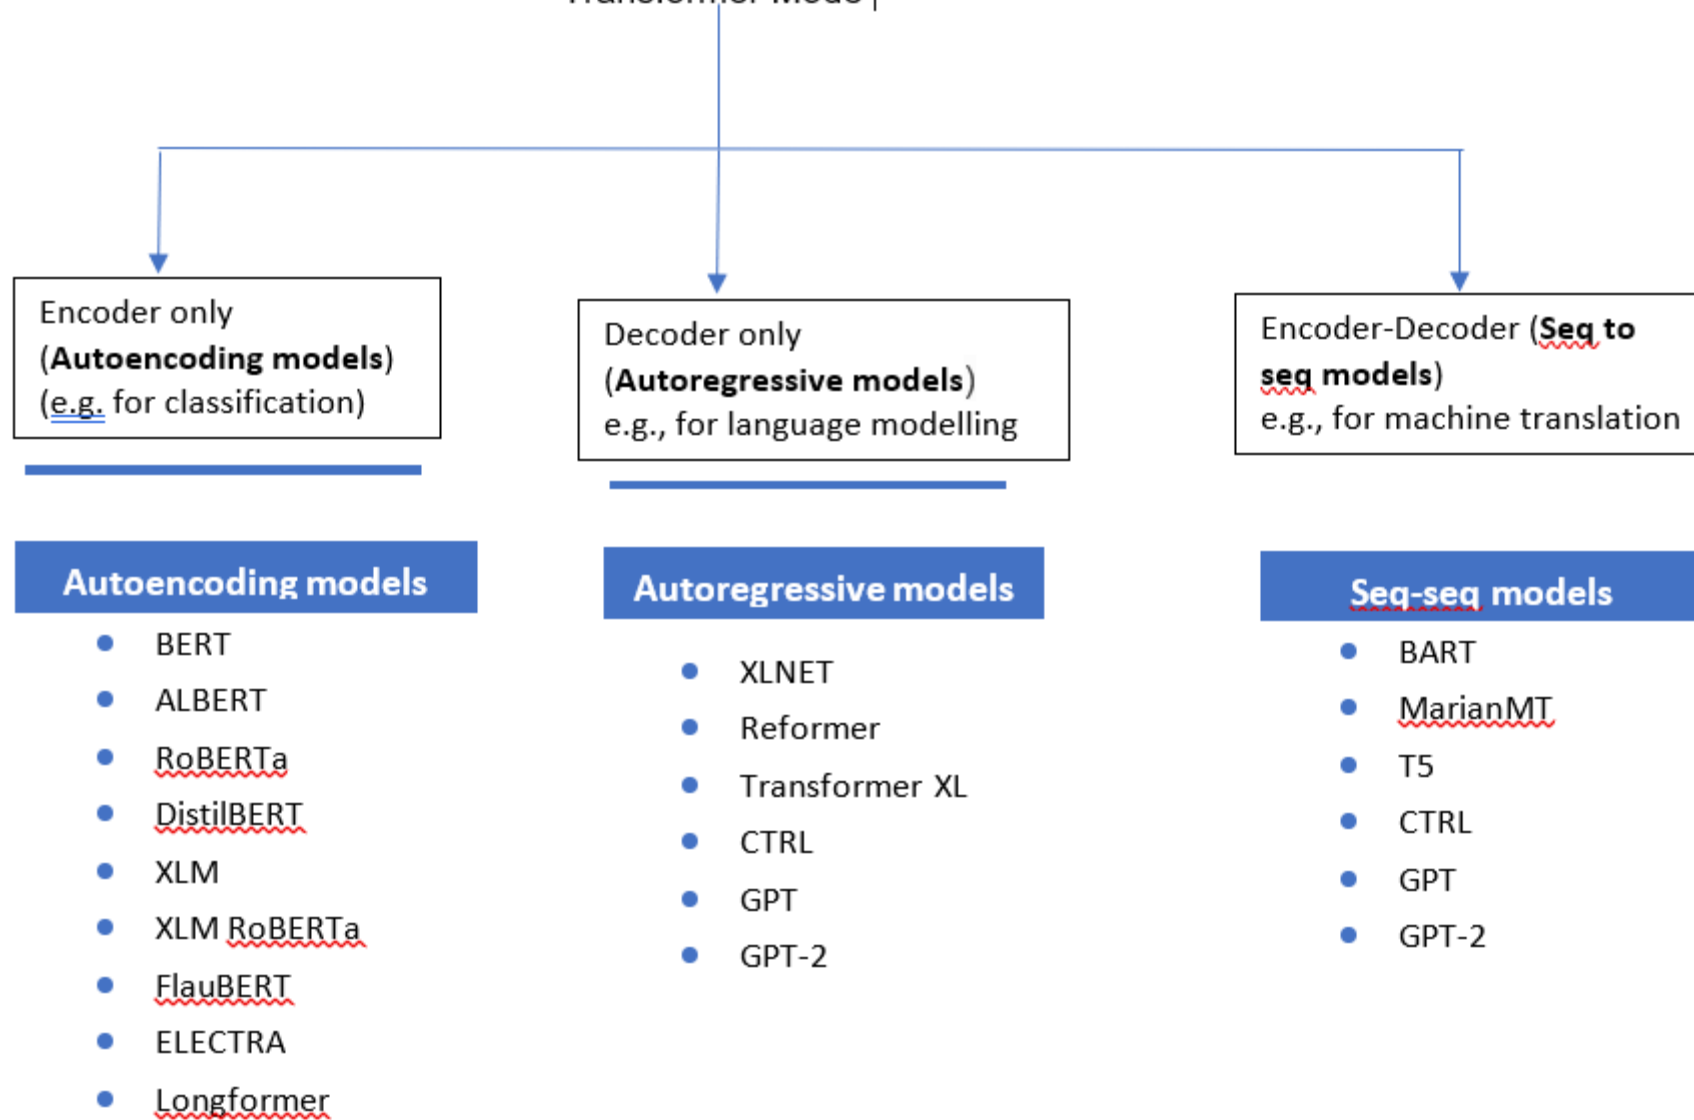
5



The self attention layers in the decoder operate in a slightly different way than the one in the encoder:

In the decoder, the self-attention layer is only allowed to attend to earlier positions in the output sequence. This is done by masking future positions (setting them to  $-\infty$ ) before the softmax step in the self-attention calculation.

## Transformer Mode |



# Deep Learning Models with attention

## Natural Language Processing

BART  
(BERT)  
RoBERTa  
DistilBERT  
Generative Pre-trained Transformer (GPT)  
GPT-2 GPT-3  
Transformer-XL  
XLNet combines BERT and Transformer-XL

## Object Detection

Deep Recurrent Attentive Writer (DRAW)

## Transformers for vision tasks

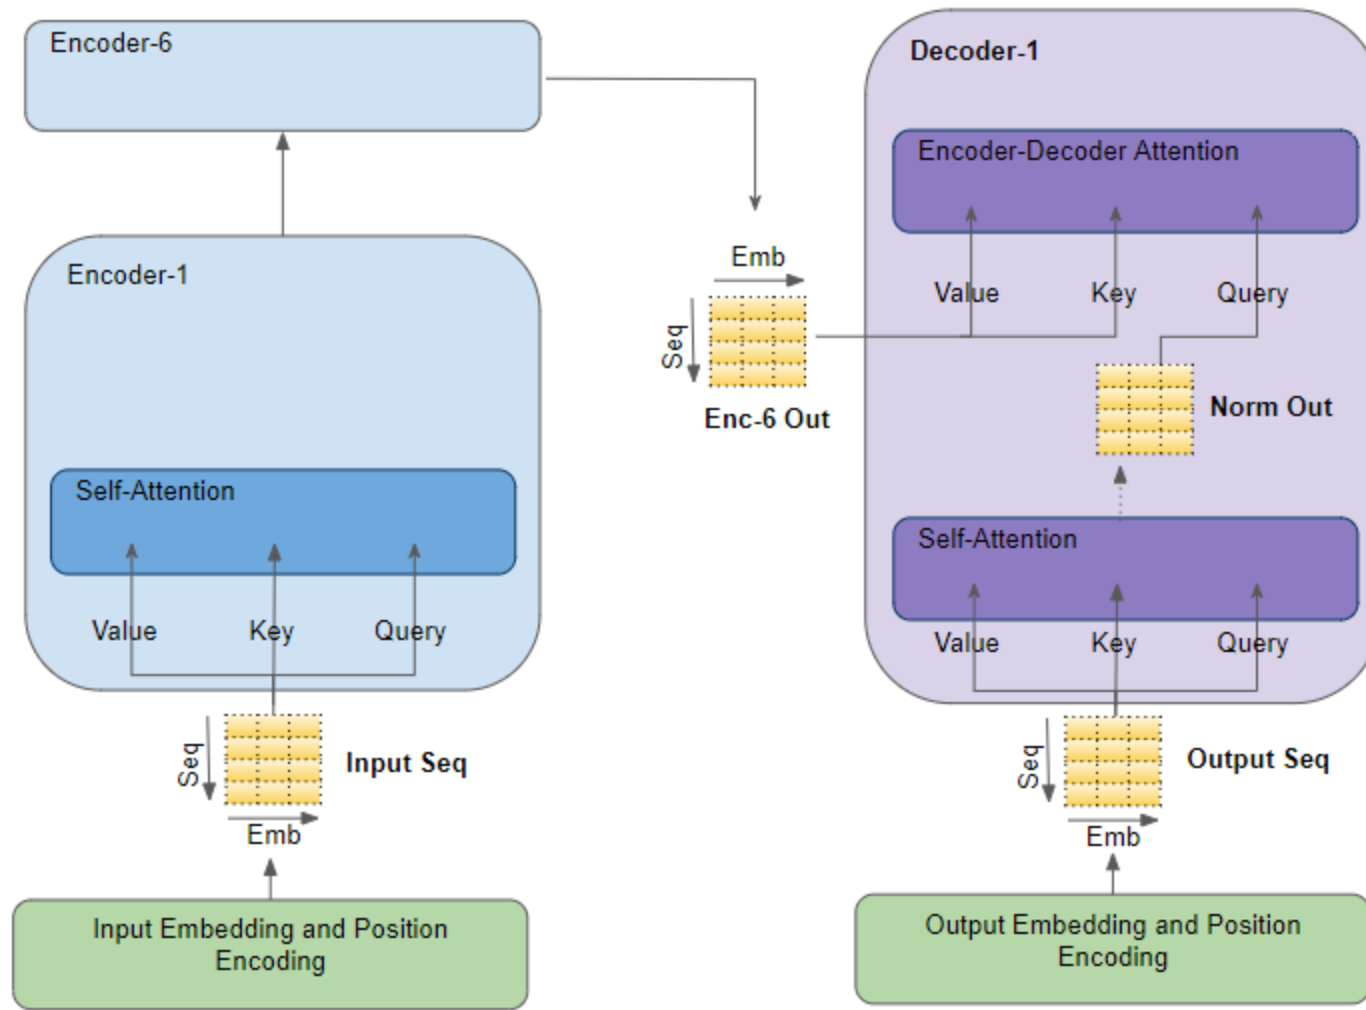
Vision Transformer (ViT)  
Detection Transformer (DETR)

## Multimodel

- videoBERT
- visualBERT

## Image Generation

Generative Adversarial Networks (SAGANs)



Thank You

# Namah Shivaya